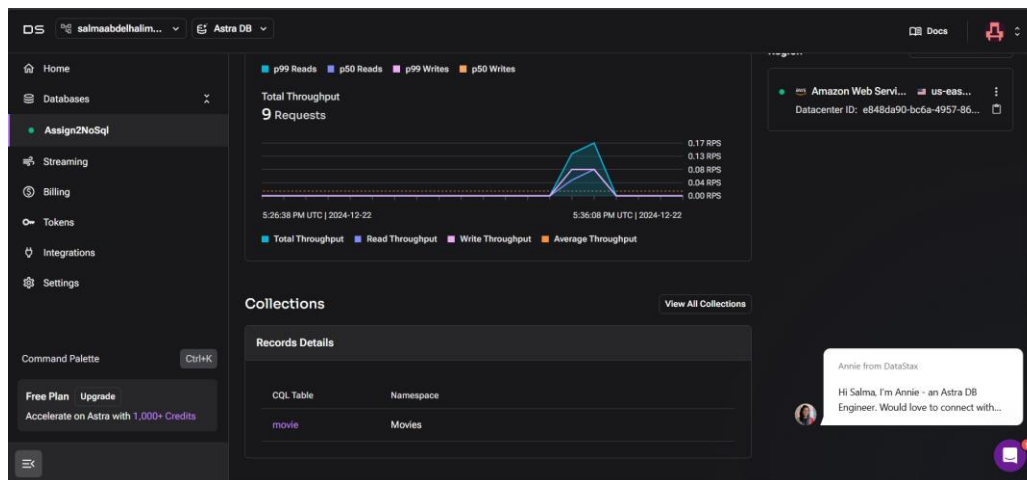
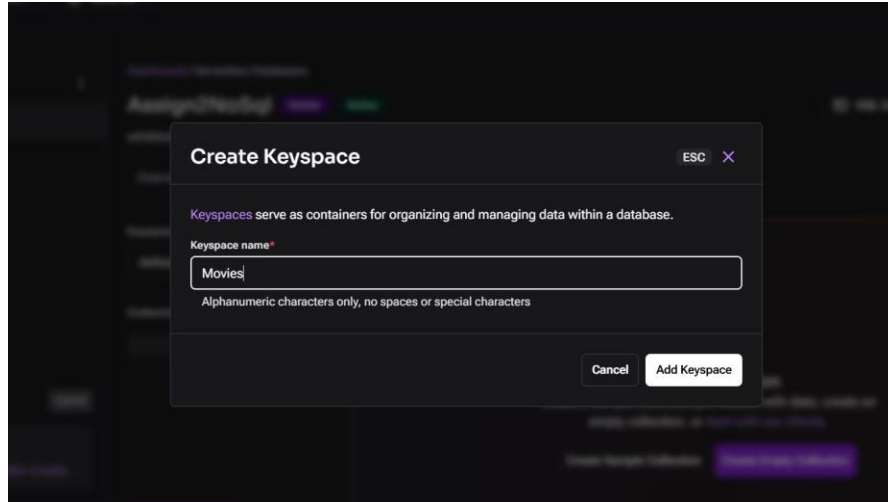


Assignment 2: Cassandra DB

1. Create a Keyspace:



```
Connected as salmaabdelhalim005zidan@gmail.com.
Connected to cndb at cassandra.ingress:9042.
[cqlsh 6.8.0 | Cassandra 4.0.0.6816 | CQL spec 3.4.5 | Native protocol v4 | TLS]
Use HELP for help.
token@cqlsh> describe keyspaces;

data_endpoint_auth  default_keyspace  system_traces  system_virtual_schema
system_auth         datastax_sla      "Movies"
system_schema       system             system_views

token@cqlsh> |
```

2. Create a column-family:

```
CREATE TABLE Movie (  
  id int PRIMARY KEY,  
  name text,  
  movie_cast map<text, text>,  
  movie_poster blob,  
  created_at timestamp);
```

```
token@cqlsh:Movies> CREATE TABLE Movie (  
    ...     id int PRIMARY KEY,  
    ...     name text,  
    ...     movie_cast map<text, text>,  
    ...     movie_poster blob,  
    ...     created_at timestamp  
    ... );
```

3. Check the schema:

```
token@cqlsh:Movies> describe Movie;  
  
CREATE TABLE "Movies".movie (  
  id int PRIMARY KEY,  
  created_at timestamp,  
  movie_cast map<text, text>,  
  movie_poster blob,  
  name text  
  WITH additional_write_policy = '99p'  
  AND bloom_filter_fp_chance = 0.01  
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
  AND comment = ''  
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.UnifiedCompactionStrategy'}  
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}  
  AND crc_check_chance = 1.0  
  AND default_time_to_live = 604800  
  AND gc_grace_seconds = 864000  
  AND max_index_interval = 2048  
  AND memtable_flush_period_in_ms = 0  
  AND min_index_interval = 128  
  AND read_repair = 'BLOCKING'  
  AND speculative_retry = '99p';  
  
token@cqlsh:Movies> █
```

4. Populate the Movie table with 3 real movies:

```
INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at)
VALUES (
  1,
  'Creed 3',
  {'director': 'Michael B. Jordan', 'actors': 'Michael B. Jordan, Tessa Thompson',
'music': 'Joseph Shirley'},
  null,
  toTimestamp(now()))
USING TTL 604800;

INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at)
VALUES ( 2, 'Forrest Gump',
  {'director': 'Robert Zemeckis', 'actors': 'Tom Hanks, Robin Wright', 'music': 'Alan
Silvestri'},
  null, toTimestamp(now()) );

INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at)
VALUES ( 3, 'Inception',
  {'director': 'Christopher Nolan', 'actors': 'Leonardo DiCaprio, Joseph Gordon-Levitt',
'music': 'Hans Zimmer'}, null, toTimestamp(now()));
```

```
token@cqlsh:Movies> INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at)
... VALUES (
... 1,
... 'Creed 3',
... {'director': 'Michael B. Jordan', 'actors': 'Michael B. Jordan, Tessa Thompson', 'music': 'Joseph Shirley'},
... null,
... toTimestamp(now()))
... USING TTL 604800;
token@cqlsh:Movies>
token@cqlsh:Movies> INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at) VALUES ( 2, 'Forrest Gump',
... {'director': 'Robert Zemeckis', 'actors': 'Tom Hanks, Robin Wright', 'music': 'Alan Silvestri'},
... null, toTimestamp(now()) );
token@cqlsh:Movies>
token@cqlsh:Movies> INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at) VALUES ( 3, 'Inception',
... {'director': 'Christopher Nolan', 'actors': 'Leonardo DiCaprio, Joseph Gordon-Levitt', 'music': 'Hans Zimmer'}, null, toTimestamp(now()));
```

```
token@cqlsh:use Movies ;
token@cqlsh:Movies> Select * from Movie;

id | created_at | movie_cast | movie_poster | name
---|---|---|---|---
1 | 2024-12-17 18:42:35.659000+0000 | {'actors': 'Michael B. Jordan, Tessa Thompson', 'director': 'Michael B. Jordan', 'music': 'Joseph Shirley'} | null | Creed 3
2 | 2024-12-17 18:42:35.668000+0000 | {'actors': 'Tom Hanks, Robin Wright', 'director': 'Robert Zemeckis', 'music': 'Alan Silvestri'} | null | Forrest Gump
3 | 2024-12-17 18:42:40.845000+0000 | {'actors': 'Leonardo DiCaprio, Joseph Gordon-Levitt', 'director': 'Christopher Nolan', 'music': 'Hans Zimmer'} | null | Inception

(3 rows)
token@cqlsh:Movies> 
```

5. The Python function connects to KeySpace transforms the movie poster into a blob datatype and updates the movie-poster.

5.1. Connect to the keySpace(Movies):

connects to KeySpace:

```
from cassandra.cluster import Cluster
from cassandra.auth import PlainTextAuthProvider
import json
cloud_config= {
    'secure_connect_bundle': '/content/secure-connect-assign2nosql.zip'}
with open("/content/salmaabdelhalim005zidan@gmail.com-token (1).json") as f:
    secrets = json.load(f)
CLIENT_ID = secrets["clientId"]
CLIENT_SECRET = secrets["secret"]
auth_provider = PlainTextAuthProvider(CLIENT_ID, CLIENT_SECRET)
cluster = Cluster(cloud=cloud_config, auth_provider=auth_provider)
session = cluster.connect('Movies')
row = session.execute("select release_version from system.local").one()
if row:
    print(row[0])
else:
    print("An error occurred.")
```



```
from cassandra.cluster import Cluster
from cassandra.auth import PlainTextAuthProvider
import json

cloud_config= {
    'secure_connect_bundle': '/content/secure-connect-assign2nosql.zip'
}
with open("/content/salmaabdelhalim005zidan@gmail.com-token (1).json") as f:
    secrets = json.load(f)

CLIENT_ID = secrets["clientId"]
CLIENT_SECRET = secrets["secret"]

auth_provider = PlainTextAuthProvider(CLIENT_ID, CLIENT_SECRET)
cluster = Cluster(cloud=cloud_config, auth_provider=auth_provider)
session = cluster.connect('Movies')

row = session.execute("select release_version from system.local").one()
if row:
    print(row[0])
else:
    print("An error occurred.")
```

WARNING:cassandra.cluster:Downgrading core protocol version from 66 to 65 for e848da90-bc6a-4957-8673-548a6835cb25-us-east-2.db.astra.datastax.com:29042:0669df0f-031f-383e-b452-67a967f140.11-7cde36df13c4
WARNING:cassandra.cluster:Downgrading core protocol version from 65 to 5 for e848da90-bc6a-4957-8673-548a6835cb25-us-east-2.db.astra.datastax.com:29042:0669df0f-031f-383e-b452-67a967f140.11-7cde36df13c4
WARNING:cassandra.cluster:Downgrading core protocol version from 5 to 4 for e848da90-bc6a-4957-8673-548a6835cb25-us-east-2.db.astra.datastax.com:29042:0669df0f-031f-383e-b452-67a967f140.11-7cde36df13c4

5.2. updates the movie-poster column:

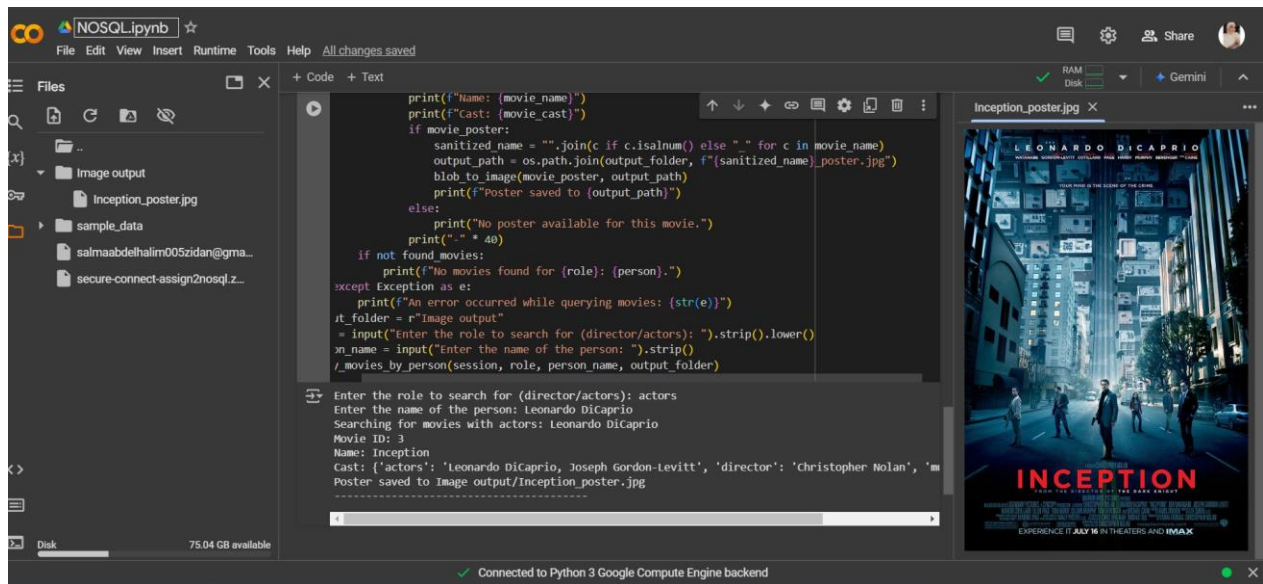
updates the movie-poster column:

```
import os
def image_to_blob(image_path):
    """
    Reads an image and converts it to a binary blob.
    """
    with open(image_path, "rb") as file:
        return file.read()
def update_movie_posters(session, folder_path):
    try:
        poster_files = [f for f in os.listdir(folder_path) if f.endswith(('jpg', 'jpeg', 'png'))]
        rows = session.execute("SELECT Id, name FROM Movie").all()
        if len(poster_files) < len(rows):
            print(f"Not enough posters to match with the movies.")
            Return
        for (movie_id, movie_name), poster_file in zip(rows, poster_files):
            poster_path = os.path.join(folder_path, poster_file)
            poster_blob = image_to_blob(poster_path)
            query = "UPDATE Movie SET movie_poster = %s WHERE Id = %s"
            session.execute(query, (poster_blob, movie_id))
            print(f"Successfully updated the poster for {movie_name} with {poster_file}")
    except Exception as e:
        print(f"An error occurred while updating posters: {str(e)}")
folder_path = r"C:\\Users\\lenovo\\Desktop\\Posters"
update_movie_posters(session, folder_path)
```


6. python function to query the movies given a certain director or actor:

```
import os
def blob_to_image(blob_data, output_path):
    with open(output_path, "wb") as file:
        file.write(blob_data)
def query_movies_by_person(session, role, person, output_folder):
    try:
        print(f"Searching for movies with {role}: {person}")
        query = "SELECT * FROM Movie ALLOW FILTERING;"
        rows = session.execute(query)
        os.makedirs(output_folder, exist_ok=True)
        found_movies = False
        for row in rows:
            movie_id = row.id
            movie_name = row.name
            movie_cast = row.movie_cast
            movie_poster = row.movie_poster
            if role in movie_cast and person in movie_cast[role]:
                found_movies = True
                print(f"Movie ID: {movie_id}")
                print(f"Name: {movie_name}")
                print(f"Cast: {movie_cast}")
                if movie_poster:
                    sanitized_name = "".join(c if c.isalnum() else "_" for c in movie_name)
                    output_path = os.path.join(output_folder, f"{sanitized_name}_poster.jpg")
                    blob_to_image(movie_poster, output_path)
                    print(f"Poster saved to {output_path}")
                else:
                    print("No poster available for this movie.")
                    print("-" * 40)
            if not found_movies:
                print(f"No movies found for {role}: {person}.")
        except Exception as e:
            print(f"An error occurred while querying movies: {str(e)}")
    output_folder = r"Image output"
    role = input("Enter the role to search for (director/actors): ").strip().lower()
    person_name = input("Enter the name of the person: ").strip()
    query_movies_by_person(session, role, person_name, output_folder)
```

```
Connected as faridahamid2004@gmail.com.
Connected to cndb at cassandra.ingress:9042.
[cqlsh 6.8.0 | Cassandra 4.0.0.6816 | CQL spec 3.4.5 | Native protocol v4 | TLS]
Use HELP for help.
Region: us-east1
token@cqlsh> use "Movies";
token@cqlsh:Movies> select * from Movie where movie_cast CONTAINS 'Tom Hanks';
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
token@cqlsh:Movies> █
```



7. Update a certain movie actors list to add another actor to the list.

UPDATE Movie SET movie_cast = movie_cast + {'actors': 'Michael B. Jordan, Tessa Thompson, Jonathan Majors'} WHERE Id = 1;

```
token@cqlsh> use "Movies";
token@cqlsh:Movies> UPDATE Movie
... SET movie_cast = movie_cast + {'actors': 'Michael B. Jordan, Tessa Thompson, Jonathan Majors'}
... WHERE Id = 1;
token@cqlsh:Movies> select Id, name, movie_cast, created_at from movie;
```

Id	name	movie_cast	created_at
1	Creed 3	{'actors': 'Michael B. Jordan, Tessa Thompson, Jonathan Majors', 'director': 'Michael B. Jordan', 'music': 'Joseph Shirley'}	2024-12-18 07:50:40.775000+0000
2	Forrest Gump	{'actors': 'Tom Hanks, Robin Wright', 'director': 'Robert Zemeckis', 'music': 'Alan Silvestri'}	2024-12-18 07:50:40.782000+0000
3	Inception	{'actors': 'Leonardo DiCaprio, Joseph Gordon-Levitt', 'director': 'Christopher Nolan', 'music': 'Hans Zimmer'}	2024-12-18 07:50:42.266000+0000

(3 rows)
token@cqlsh:Movies>

8. Update the first row TTL to 3 seconds:

INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at) VALUES (1, 'Creed 3', {'director': 'Michael B. Jordan', 'actors': 'Michael B. Jordan, Tessa Thompson, Jonathan Majors', 'music': 'Joseph Shirley'}, 0x48656c6f2c20576f726c6421 ,toTimestamp(now())) USING TTL 3;

```
token@cqlsh:Movies> INSERT INTO Movie (Id, name, movie_cast, movie_poster, created_at) VALUES ( 1, 'Creed 3', {'director': 'Michael B. Jordan', 'actors': 'Michael B. Jordan, Tessa Thompson, Jonathan Majors', 'music': 'Joseph Shirley'}, 0x48656c6f2c20576f726c6421 ,toTimestamp(now())) USING TTL 3;
token@cqlsh:Movies> select Id, name, movie_cast, created_at from movie;
```

Id	name	movie_cast	created_at
2	Forrest Gump	{'actors': 'Tom Hanks, Robin Wright', 'director': 'Robert Zemeckis', 'music': 'Alan Silvestri'}	2024-12-18 07:50:40.782000+0000
3	Inception	{'actors': 'Leonardo DiCaprio, Joseph Gordon-Levitt', 'director': 'Christopher Nolan', 'music': 'Hans Zimmer'}	2024-12-18 07:50:42.266000+0000

(2 rows)
token@cqlsh:Movies>